

EEE554 – Probability and Random Processes

Module 3: Conditional Probability and Bayes' Rule

Nicolò Michelusi

Associate Professor

School of Electrical, Computer and Energy Engineering

Arizona State University

From Probability to Conditional Probability

So far, probabilities have been computed using a probability space

$$(\Omega, \mathcal{F}, P).$$

For an event

$$A \in \mathcal{F},$$

we compute

$$P(A).$$

This represents our probability of A given the information currently available. But what happens when we acquire new information?

Why Conditional Probability?

In many applications, we accumulate *additional information*:

- Probability of packet loss given network congestion.
- Probability of a communication error given the transmitted symbol.
- Probability of disease given a positive medical test.
- Probability of target detection given the received signal.

Conditioning incorporates the available information into the probability calculation.

Conditional Probability

Suppose we learn that an event B has occurred.

Outcomes outside B are no longer possible.


Question:

How should this new information change the probability of A ?

Knowing that B occurred restricts our attention to B .

Within this restricted sample space, the outcomes for which A also occurs are

$$A \cap B.$$



$$P(A|B)$$

Therefore, we renormalize probability within B :

$$P(A|B) = \frac{P(A \cap B)}{P(B)}.$$

$$B = (A \cap B) \cup (A^c \cap B)$$

disjoint

This naturally leads to the following definition.


Definition 1. For $A, B \in \mathcal{F}$ with

$$P(B) > 0,$$

the conditional probability of A given B is

$$P(A|B) = \frac{P(A \cap B)}{P(B)}.$$

$$P(A^c | B) = \frac{P(A^c \cap B)}{P(B)}$$



$$P(B) = P(A \cap B) + P(A^c \cap B)$$

↓


$$1 = P(A|B) + P(A^c|B)$$

$$P(B^c|B) = \frac{P(B^c \cap B)}{P(B)} = \frac{P(\emptyset)}{P(B)} = 0$$

Properties

If $A \subseteq B$, then $A \cap B = A$

$$P(A|B) = \frac{P(A)}{P(B)} \geq P(A).$$




Knowing that B occurred makes A more likely.

If $B \subseteq A$, then

$$P(A|B) = 1.$$


Once B occurs, A is certain.



$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B)}{P(B)} = 1$$

If $A \cap B = \emptyset$, then

$$P(A|B) = 0.$$



Knowing that B occurred makes A impossible, since no outcome lies in both A and B .

Conditioning Defines a New Probability Measure

For fixed B , with $P(B) > 0$,

$$P(\cdot|B)$$

is itself a probability measure.

Thus,

$$\boxed{(\Omega, \mathcal{F}, P) \longrightarrow (\Omega, \mathcal{F}, P(\cdot|B)).}$$

Conditioning updates our probability model after learning that B occurred.

Bayes' Theorem

If we know $P(B|A)$, can we find $P(A|B)$?

Observe that

$$P(A \cap B) = P(A|B)P(B) = P(B|A)P(A).$$

Hence,

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

$$P(A|B) = \frac{P(A \cap B)}{P(B)} ; P(B|A) = \frac{P(B \cap A)}{P(A)}$$

Law of Total Probability

Recall that a partition

$$A_1, \dots, A_n$$

divides the sample space into disjoint events:


$$\Omega = A_1 \cup \dots \cup A_n,$$

disjoint!

Therefore,

$$B = (B \cap A_1) \cup \dots \cup (B \cap A_n).$$

$$B = B \cap \Omega = B \cap (A_1 \cup \dots \cup A_n)$$



$$A_i \cap A_j = \emptyset \quad \forall i \neq j$$

Since these events are disjoint,

$$P(B) = \sum_{i=1}^n P(B \cap A_i).$$

Using

$$P(B \cap A_i) = P(B|A_i)P(A_i),$$

we obtain

$$P(B) = \sum_{i=1}^n P(B|A_i)P(A_i).$$

Special case: $\Omega = A \cup A^c \Rightarrow P(B) = P(B|A) \cdot P(A) + P(B|A^c) \cdot P(A^c)$

Bayes' Theorem with a Partition

Combining Bayes' theorem with the law of total probability,

$$P(A_i|B) = \frac{P(B|A_i)P(A_i)}{\sum_{j=1}^n P(B|A_j)P(A_j)}$$

$$P(A_i|B) = \frac{P(B|A_i) \cdot P(A_i)}{P(B)}$$

Interpretation of Bayes' Rule

$$P(A_i|B) = \frac{P(B|A_i)P(A_i)}{P(B)}.$$

Handwritten blue annotations:
An arrow points from "bit 1 tx" to $P(A_i|B)$.
An arrow points from "bit 1 rx" to $P(B)$.

$$\underbrace{P(A_i)}_{\text{Prior}} \quad \underbrace{P(B|A_i)}_{\text{Likelihood}} \quad \underbrace{P(A_i|B)}_{\text{Posterior}}$$

where


$$P(B) = \sum_j P(B|A_j)P(A_j)$$

is the probability of observing the evidence B .


Prior belief + new observation $\longrightarrow$ updated belief.

$P(A) \longrightarrow \text{observe } B \text{ happened}$
 $\downarrow$
 $P(A|B)$

Engineering Example: Detection



An unknown symbol $S \in \{1, \dots, N\}$ is transmitted. At the receiver, we observe $Y \in \{1, \dots, N\}$. The goal is to estimate (infer) the transmitted symbol. We call this estimate $\hat{S}$.



The channel is characterized by

$$P(Y = y|S = s).$$

For a fixed observation $Y = y$, $P(Y = y|S = s)$, viewed as a function of s , is called the **likelihood**.

It quantifies how likely the observation $Y = y$ is when symbol $S = s$ is transmitted.

Posterior probability

After observing $Y = y$, what is the probability that symbol s was transmitted?

Answer:

$$P(S = s|Y = y),$$

called the **posterior probability**.

Bayes' theorem computes the posterior from the likelihood:

$$P(S = s | Y = y) = \frac{P(Y = y | S = s) P(S = s)}{\sum_{i=1}^N P(Y = y | S = i) P(S = i)}.$$

Thus, posterior probabilities can be computed from

- prior probabilities,
- likelihoods,
- the law of total probability.


$$\frac{P(Y=y|S=s) \cdot P(S=s)}{P(Y=y)}$$

Maximum Likelihood (ML) Detection

Suppose the prior probabilities are unknown. Choose

$$\hat{s}_{ML} = \arg \max_s P(Y = y | S = s).$$

This is called a *Maximum Likelihood detector*. ML ignores prior probabilities.

Maximum A Posteriori (MAP) Detection

If the prior probabilities are available,

$$\hat{s}_{MAP} = \arg \max_s P(S = s | Y = y).$$

This is called a *Maximum A Posteriori* detector. MAP exploits prior information.

$$= \arg \max_s \frac{p(Y=y | S=s) \cdot p(S=s)}{p(Y=y)}$$


Using Bayes,

$$P(Y=y \wedge S=s)$$

$$\hat{s}_{MAP} = \arg \max_s P(Y = y|S = s)P(S = s),$$

since the denominator does not depend on s .

Example: Binary Symmetric Channel



Source:

$$P(S = 0) = 0.9, \quad P(S = 1) = 0.1.$$

Channel:

$$P(Y \neq s | S = s) = \epsilon, \text{ where } \epsilon \leq \frac{1}{2}.$$

Maximum Likelihood (ML)

Suppose $Y = 1$ is observed. The likelihoods are

$$P(Y = 1|S = 0) = \varepsilon,$$

$$P(Y = 1|S = 1) = 1 - \varepsilon.$$

Since $1 - \varepsilon \geq \varepsilon$, the ML estimator is $\hat{S}_{\text{ML}}(1) = 1$. Similarly, if $Y = 0$ is observed, then $\hat{S}_{\text{ML}}(0) = 0$. Overall:

$$\boxed{\hat{S}_{\text{ML}}(Y) = Y.}$$

Maximum A Posteriori (MAP)

Using Bayes' theorem,

$$P(S = s|Y = 1) \propto P(Y = 1|S = s)P(S = s).$$

Hence,

$$P(S = 0|Y = 1) \propto 0.9\varepsilon,$$

$$P(S = 1|Y = 1) \propto 0.1(1 - \varepsilon).$$

Therefore,

$$\hat{S}_{\text{MAP}}(1) = \begin{cases} 0, & 0.9\varepsilon \geq 0.1(1 - \varepsilon), \\ 1, & 0.9\varepsilon < 0.1(1 - \varepsilon). \end{cases}$$

Similarly:

$$\hat{S}_{\text{MAP}}(0) = \begin{cases} 0, & 0.9(1 - \varepsilon) \geq 0.1\varepsilon, \\ 1, & 0.9(1 - \varepsilon) < 0.1\varepsilon \end{cases}$$

In other words, the MAP detector looks like:

if $0 \leq \varepsilon \leq 0.1$, then $\hat{S}_{\text{MAP}}(Y) = Y$; $\Rightarrow \hat{S}_{\text{ML}}(Y)$

if $0.5 \geq \varepsilon > 0.1$, then $\hat{S}_{\text{MAP}}(Y) = 0$.

Probability of Error

The probability of error is

$$P_e = P(\hat{S} \neq S).$$

For a BSC with $\varepsilon \leq \frac{1}{2}$, the ML detector is

$$\hat{S}_{\text{ML}}(Y) = Y.$$

Hence,

$$P_e^{\text{ML}} = P(\hat{S}_{\text{ML}} \neq S) = P(Y \neq S)$$

$$= P(\underbrace{Y \neq S}_{Y \neq 0 \equiv Y=1} | S=0)P(S=0) + P(\underbrace{Y \neq S}_{Y \neq 1 \equiv Y=0} | S=1)P(S=1) = \varepsilon.$$

$$Y \neq 0 \equiv Y=1$$

$$Y \neq 1 \equiv Y=0$$

$$= \varepsilon \cdot P(S=0) + \varepsilon \cdot P(S=1) = \varepsilon$$

The MAP rule is

$$\hat{S}_{\text{MAP}}(Y) = \begin{cases} 0, & \varepsilon > \frac{1}{10}, \\ Y, & \varepsilon \leq \frac{1}{10}. \end{cases}$$

Therefore,

$$P_e^{\text{MAP}} = \begin{cases} \varepsilon, & 0 \leq \varepsilon < \frac{1}{10}, \\ 0.1, & \frac{1}{10} \leq \varepsilon < \frac{1}{2}. \end{cases}$$

$$0.1 < \varepsilon \leq 0.5$$

$$P(\hat{S}_{\text{MAP}} \neq S) = P(S \neq 0) = P(S=1) = 0.1$$

Remark.

$$P_e^{\text{MAP}} \leq P_e^{\text{ML}}$$

with strict inequality for

$$\frac{1}{10} < \varepsilon < \frac{1}{2}.$$

Prior information improves detection performance.